

Emerging Safety Attack and Defense in Federated Instruction Tuning of Large Language Models

Authors: Rui Ye, Jingyi Chai, Xiangrui Liu, Yaodong
Yang, Yanfeng Wang, Siheng Chen

Conference: ICLR 2025

Presented by: Maha Zainab

Learning Objectives

- ▶ Understand the concept of Federated Instruction Tuning (FedIT).
- ▶ Recognize the emerging safety risks in FedIT.
- ▶ Analyze how malicious clients can compromise safety.
- ▶ Evaluate why existing defenses fail.
- ▶ Explore the proposed post-hoc defense strategy and its effectiveness.

The Challenge of FedIT

- ▶ Federated Instruction Tuning (FedIT) allows collaborative training without data sharing.
- ▶ LLMs need instruction tuning for helpfulness + safety by training on decentralized data.
- ▶ Challenge: FedIT can be vulnerable to attacks.

The Stealthy Safety Attack


- ▶ Malicious clients generate harmful instructions and responses
- ▶ Local Models trained on unsafe data
- ▶ Global LLM loses safety alignment

Research Questions

- ▶ **Can malicious clients compromise safety in FedIT?**
 - ▶ Simple unaligned data (harmful instructions → harmful answers) from a few malicious clients can break the global model's safety alignment.
- ▶ **Why do existing federated defenses fail?**
 - ▶ Classical defenses (e.g., robust aggregation) are ineffective; malicious updates appear almost identical to benign ones.
- ▶ **Can we design an effective defense?**
 - ▶ A post-hoc, server-side alignment step may restore safety without harming helpfulness.

Why Existing Defenses Fail?

- ▶ Clients locally fine-tune LLMs on private data.
- ▶ Server aggregates updates → global LLM.
- ▶ Goal: leverage diverse data, preserve privacy, and ensure safety.



Why Existing Defenses Fail

- ▶ Classical defenses rely on model-level filtering.
- ▶ Benign & malicious updates look similar (both 'follow instructions').
- ▶ Safety drops by up to 70%.
- ▶ Existing defenses improve safety by $\leq 4\%$.


Proposed Post-Hoc Defense

- ▶ Server runs extra fine-tuning after aggregation.
- ▶ Generates defense dataset:
 - ▶ - Normal instructions → helpful answers
 - ▶ - Harmful instructions → harmless refusals
- ▶ Fine-tunes global LLM on this dataset.

Proposed Post-Hoc Defense

Automated Data Generation

- LLM generates harmful/normal instructions.
- LLM generates harmless responses for harmful instructions, and helpful responses for normal instructions.






Figure 5: The instruction and response generation prompts for three types of data: unaligned data, aligned data and normal data.

Experimental Setup

- ▶ Base model: Llama2-7B + LoRA.
- ▶ Datasets: LMSYS-Chat, WildChat, Beavertails, MaliciousGen.
- ▶ 10 clients (7 benign, 3 malicious).
- ▶ Benchmarks: AdvBench (safety), MT-Bench (helpfulness).

Results: Beavertails and Two Benign Datasets

Benign Dataset Evaluation Metric $\uparrow$	LMSYS-Chat				WildChat			
	Rule	MD-Judge	RM	MT-1	Rule	MD-Judge	RM	MT-1
FedAvg (No Attack)	82.88	66.15	-1.72	4.19	79.04	43.27	-1.63	4.75
FedAvg	49.81	25.96	-2.97	4.14	38.65	12.31	-2.73	4.54
Median	48.65	23.85	-3.10	3.88	41.35	10.58	-2.80	4.74
Trimmedmean	45.96	26.35	-3.05	4.20	41.35	14.04	-2.84	4.43
Krum	55.38	27.88	-2.88	4.16	40.00	9.42	-2.48	4.55
DnC	55.96	25.38	-2.90	4.00	41.15	7.12	-2.63	4.41
FoolsGold	46.92	25.00	-3.05	3.95	37.50	10.96	-2.79	4.55
Residual	47.50	23.65	-2.98	4.04	37.50	10.77	-2.86	4.54
Ours: Level 1	68.65	44.23	-2.31	4.11	57.31	17.50	-2.26	4.85
Ours: Level 2	77.31	84.23	-0.99	4.23	82.12	82.12	-1.08	4.33
Ours: Level 3	62.69	72.88	-1.65	3.73	51.54	57.69	-1.90	4.39

Results: Malicious-Gen and Two Benign Datasets

Benign Dataset Evaluation Metric $\uparrow$	LMSYS-Chat				WildChat			
	Rule	MD-Judge	RM	MT-1	Rule	MD-Judge	RM	MT-1
FedAvg (No Attack)	82.88	66.15	-1.72	4.19	79.04	43.27	-1.63	4.75
FedAvg	43.27	11.35	-3.62	4.19	30.58	5.78	-3.03	4.40
Median	48.27	13.65	-3.43	3.95	40.00	10.19	-3.02	4.10
Trimmedmean	41.92	9.62	-3.51	3.71	31.92	5.96	-3.13	4.09
Krum	50.38	16.73	-3.23	4.14	39.04	7.89	-2.99	4.55
DnC	49.04	12.12	-3.40	4.14	45.58	9.04	-2.90	4.49
FoolsGold	41.54	12.12	-3.45	3.85	30.78	6.35	-3.03	4.14
Residual	44.23	10.19	-3.52	3.80	31.54	6.15	-3.00	4.14
Ours: Level 1	71.15	34.32	-2.68	4.19	50.38	13.27	-2.18	4.61
Ours: Level 2	78.08	83.08	-0.96	4.18	77.12	72.50	-1.49	4.13
Ours: Level 3	75.96	72.69	-1.56	3.89	58.08	62.12	-1.70	4.33

Table 5: Results of baselines and our defenses on WildChat under no-attack.

Evaluation Metric $\uparrow$	Rule	MD-Judge	RM	MT-1
FedAvg	79.04	43.27	-1.63	4.75
Median	79.81	44.23	-1.50	4.70
Trimmedmean	80.58	44.04	-1.65	4.36
Krum	78.08	45.19	-1.53	4.54
DnC	77.50	40.77	-1.75	4.58
FoolsGold	80.78	46.15	-1.59	4.36
Residual	78.08	40.00	-1.69	4.49
Ours: Level 1	76.35	41.35	-1.67	4.89
Ours: Level 2	82.31	74.62	-1.33	4.24
Ours: Level 3	88.08	69.62	-1.16	4.65

Plug and Play

Metrics ↑	+ Ours	FedAvg	Median	Trimmed.	Krum	DnC	FoolsGold	Residual
Rule	X	49.81	48.65	45.96	55.38	55.96	46.92	47.50
	✓	77.31	77.88	79.42	79.42	80.00	81.35	78.08
MD-J	X	25.96	23.85	26.35	27.88	25.38	25.00	23.65
	✓	84.23	86.35	84.04	82.31	84.42	88.08	86.92
RM	X	-2.97	-3.10	-3.05	-2.88	-2.90	-3.05	-2.98
	✓	-1.00	-0.92	-1.10	-1.02	-1.07	-0.98	-0.94
MT-1	X	4.14	3.88	4.20	4.16	4.00	3.95	4.04
	✓	4.14	4.06	3.95	3.88	4.01	3.94	4.29

Scalability

Client Number Evaluation Metric $\uparrow$	K=50				K=100			
	Rule	MD-Judge	RM	MT-1	Rule	MD-Judge	RM	MT-1
FedAvg (No Attack)	77.12	55.96	-1.76	4.20	79.23	54.62	-1.90	4.23
FedAvg	40.58	11.35	-3.58	3.86	37.31	9.42	-3.58	3.93
Krum	45.00	10.77	-3.56	4.09	45.19	14.04	-3.40	4.28
DnC	46.92	12.88	-3.66	4.19	46.54	15.19	-3.48	4.34
Ours	81.73	80.77	-1.08	4.34	79.23	82.12	-0.95	4.24

Table 6: Results of baselines and our defenses on multi-domain datasets mixed with 250 samples of LMSYS-Chat and 250 samples of CodeAlpaca.

Evaluation Metric $\uparrow$	Rule	MD-Judge	RM	MT-1	HumanEval pass@1
FedAvg (No Attack)	60.00	42.12	-2.15	4.08	17.07
FedAvg	35.19	7.50	-3.77	3.86	14.63
Krum	39.42	12.12	-3.51	4.13	17.68
DnC	39.04	11.73	-3.71	4.41	18.29
Ours: Level 1	55.96	25.77	-2.94	4.50	15.24
Ours: Level 2	76.73	87.88	-0.79	4.11	17.68

Key Insights

- ▶ Safety in FedIT is highly vulnerable.
- ▶ Malicious clients can attack cheaply & stealthily.
- ▶ Post-hoc server-side defense is effective and scalable

Conclusion

- ▶ First paper to expose safety risks in FedIT.
- ▶ Novel safety attack + effective defense.
- ▶ Roadmap: Train collaboratively via FedIT → Apply post-training safety alignment.
- ▶ Ensures responsible deployment of LLMs.

Improvement 1: Expanding Attack Scenarios

- ▶ What can be improved?
 - ▶ Current study only considers data poisoning attacks.
 - ▶ Real attackers may use multiple strategies together.
- ▶ How will we improve it?
 - ▶ Simulate hybrid attacks by combining data poisoning and model poisoning.
 - ▶ Test defenses against adaptive attackers who change strategies.
 - ▶ Measure safety and helpfulness impact under these stronger threats.

Improvement 2: Enhancing Evaluation

- ▶ What can be improved?
 - ▶ Current evaluation uses limited benchmarks.
 - ▶ Real-world misuse often happens in multi-turn or long conversations.
- ▶ How will we improve it?
 - ▶ Add multi-turn dialogue benchmarks to test sustained safety.
 - ▶ Run dynamic stress tests like contextual jailbreaks and long-context attacks.
 - ▶ Compare results across broader alignment datasets for robustness.

Improvement 3: Evaluating Attacking Scenarios

- ▶ What can be improved
 - ▶ The safety evaluation capabilities of LLMs
 - ▶ In which type of attacks do the LLMs struggle the most?
- ▶ How we will evaluate it
 - ▶ Through Single-turn and multi-turn interventions
 - ▶ Evaluating the FIT LLMs against different attack types



Thank You
Questions?